MioHash: An Memory-Efficient and I/O-Optimized Hashing Index on Hybrid PMem-DRAM Memories

Zixiang Yu*, Zhirong Shen*, Qingjie Zeng*, Yijie Zhong*, Jiwu Shu*†,
*Xiamen University, †Tsinghua University
yuzixiang23@foxmail.com, shenzr@xmu.edu.cn, qjzeng503@gmail.com, yijiezhong@stu.xmu.edu.cn, shujw@tsinghua.edu.cn

I. SUMMARY OF IMPLEMENTATIONS

We implement our proposed MioHash with C++. We elaborate on the dependent software, representative hash indexes, and testing tools as follows.

Dependent softwares: the implementation of MioHash relies on a suite of software, including `ndctl`, `ipmctl`, `Intel PMWatch`, and `PMDK`, which can be reached via the following URLs.

```
$ wget https://github.com/pmem/ndctl
$ wget https://github.com/intel/ipmctl
$ wget https://github.com/intel/intel-pmwatch
$ wget https://github.com/pmem/pmdk
```

Representative hash indexes: we use four representative hash indexes for comparison, namely CCEH [1], Dash [2], level hashing [3], clevel hashing [4] ,Viper [5] and Plush [6], which can be downloaded via the following URLs. Note that we used the PMDK version of the level hashing implemented by Dash and add uniqueness check to CCEH for index correctness.

```
$ git clone  https://github.com/DICL/CCEH.git
$ git clone https://github.com/pmdk/dash.git
$ git clone https://github.com/chenzhangyu/
Clevel-Hashing.git
$ git clone https://github.com/tum-db/Plush.git
$ git clone https://github.com/hpides/viper.git
```

Testing tools: we employ YCSB in our evaluation. It can be reached via the following URL.

```
$ wget https://github.com/brianfrankcooper
/YCSB/releases/download/0.17.0/ycsb-0.17.0
.tar.gz
```

II. EVALUATION DETAILS

Hardware configurations: We conduct extensive experiments on a single machine equipped with two 2.10 GHz Intel Golden 5318Y CPUs (with 24 cores in total), 128 GB of DRAM memory and four Optane PMem DIMMs of 200 series (128 GB per DIMM and 512 GB in total). The Optane PMem DIMMs are configured in interleaved AppDirect Mode, which are connected to one processor. Figure 1 shows the major hardware configurations of our testbed and Figure 2 shows detailed information about the Optane PMem used in our evaluation.

```
Architecture:              x86_64
CPU op-mode(s):            32-bit, 64-bit
Byte Order:                Little Endian
Address sizes:             46 bits physical, 57 bits virtual
CPU(s):                    48
On-line CPU(s) list:       0-47
Thread(s) per core:        1
Core(s) per socket:        24
Socket(s):                 2
NUMA node(s):              2
Vendor ID:                 GenuineIntel
CPU family:                6
Model:                     106
Model name:                Intel(R) Xeon(R) Gold 5318Y CPU @ 2.10GHz
Stepping:                  6
CPU MHz:                   800.779
CPU max MHz:               3400.0000
CPU min MHz:               800.0000
BogoMIPS:                  4200.00
Virtualization:            VT-x
L1d cache:                 2.3 MiB
L1i cache:                 1.5 MiB
L2 cache:                  60 MiB
L3 cache:                  72 MiB
```

Figure 1: Configurations of our testbed.

```
DimmID | Capacity    | LockState        | HealthState | FWVersion
=================================================================
0x0010 | 126.742 GiB | Disabled, Frozen | Healthy     | 02.02.00.1516
0x0210 | 126.742 GiB | Disabled, Frozen | Healthy     | 02.02.00.1516
0x0110 | 126.742 GiB | Disabled, Frozen | Healthy     | 02.02.00.1516
0x0310 | 126.742 GiB | Disabled, Frozen | Healthy     | 02.02.00.1516
```

Figure 2: The configurations of the Intel Optane PMem DIMMs of 200 series used in the evaluation.

We perform extensive testbed experiments, including (i) experiments to understand the property of MioHash, and (ii) experiments to understand the sensitivity of MioHash.

Experiments with YCSB: We measure the access performance of MioHash, CCEH, Dash, level hashing, clevel hashing, Viper and Plush using YCSB. We generate the workloads through YCSB and then modify their format so that they can be used for testing. We then use these workloads to test each hashing index.

The following script is an example of how to generate an insert-only workload using YCSB:

```bash
#!/bin/bash
$ ycsb load basic -P uniformworkload > uniform/ \
load.text
$ ycsb run basic -P insertworkload > uniform/ \
insert.text
```

We run the following script to modify the format of workloads.

```bash
#!/bin/bash
$ dir=./uniform/load.txt
```

```
$ sed -i '1,17d' $dir
$ sed -i 's/READ usertable user/r /' $dir
$ sed -i 's/INSERT usertable user/i /' $dir
$ sed -i 's/UPDATE usertable user/u /' $dir
$ sed -i 's/\[ field0=/ /' $dir
$ sed -i 's/\ ]/ /' $dir
```

After modifying the format, we measure the performance of all hashing indexes. Due to space limits, we present the following script to evaluate MioHash. The experiments with other hashing indexes are similar.

```bash
#!/bin/bash
dir=/mnt/pmem1-node1/pool1
loadsize=50000000
runsize=50000000
FLAG='-mavx -mavx2 -mavx512f -mavx512bw -mavx512vl'

g++ -O3 -std=c++17 -I./ -lrt -c -o src/MioHash.o \
-g src/MioHash.cpp -DINPLACE -lpmemobj -lpmem \
$FLAG -lprofiler -lunwind

g++ -O3 -std=c++17 -I./ -lrt  -o bin/run_all_ycsb \
-g src/run_all_ycsb.cpp src/MioHash.o include \
/dlock.o -lpmemobj -lpmem -lpthread  \
-DMULTITHREAD -lpqos

outdir=ycsb_all
for test in 1 2 3 4
do
        for oper in "w50r50.txt"
        do
        for num in 1 2 4 8 16 24
        do
            echo " " >> $outdir
            echo "test zipfian  "$num" " \
            $oper >> $outdir
            ycsb_dir="../ycsb/zipfian/"$oper
            load_dir="../ycsb/zipfian/load.txt"
            rm -rf $dir

            numactl -N 1 -m 1 \
            ./bin/run_all_ycsb $dir \
            $loadsize $runsize $num \
            $ycsb_dir $load_dir \
            |grep -a -E "num|Ops|Start"  \
            >> $outdir
        done
    done
done
```

Experiments on Sensitivity: We also use the impact of the P-bucket size (Exp#11) as an example to clarify how we assess the performance of MioHash. Other experiments are similar.

```bash
#!/bin/bash
dir=/mnt/pmem1-node1/pool1
num=1
loadsize=50000000
```

```bash
FLAG='-mavx -mavx2 -mavx512f -mavx512bw -mavx512vl'
outdir=./sendata
srcdir=./sensitive.cpp
for testnum in 1 2 3 4 5
do
    for sbucketsize in 64 128 256 512 1024
    do
        sed -i '30d' ./src/MioHash.h
        sed -i '30i #define S_Bucket_Size '\
        $sbucketsize'' ./src/MioHash.h
        rm -rf $dir

        g++ -O3 -std=c++17 -I./ -lrt -c \
        -o src/MioHash.o -g src/MioHash.cpp \
        -DINPLACE -lpmemobj -lpmem \
        $FLAG -lprofiler -lunwind

        g++ -O3 -std=c++17 -I./ -lrt  \
        -o bin/multi_threaded_MioHash -g \
        src/test.cpp src/MioHash.o include/dlock.o \
        -lpmemobj -lpmem -lpthread \
        -DMULTITHREAD -lpqos

        numactl -N 1 -m 1 \
        ./bin/multi_threaded_MioHash $dir $loadsize
         $num |grep -a "Ops/sec"  $outdir
    done
    sed -i '30d' ./src/MioHash.h
    sed -i '30i #define S_Bucket_Size 256' \
./src/MioHash.h
done
```

REFERENCES

[1] M. Nam, H. Cha, Y.-r. Choi, S. H. Noh, and B. Nam, "Write-optimized dynamic hashing for persistent memory," in *17th USENIX Conference on File and Storage Technologies (FAST 19)*, 2019, pp. 31–44.

[2] B. Lu, X. Hao, T. Wang, and E. Lo, "Dash: Scalable hashing on persistent memory," *Proceedings of the VLDB Endowment*, pp. 1147–1161, 2020.

[3] P. Zuo, Y. Hua, and J. Wu, "Write-optimized and high-performance hashing index scheme for persistent memory," in *13th USENIX Symposium on Operating Systems Design and Implementation (OSDI 18)*, 2018, pp. 461–476.

[4] Z. Chen, Y. Hua, B. Ding, and P. Zuo, "Lock-free concurrent level hashing for persistent memory," in *2020 USENIX Annual Technical Conference (USENIX ATC 20)*, 2020, pp. 799–812.

[5] L. Benson, H. Makait, and T. Rabl, "Viper: An Efficient Hybrid PMem-DRAM Key-Value Store," *Proceedings of the VLDB Endowment*, vol. 14, no. 9, pp. 1544–1556, 2021.

[6] L. Vogel, A. Van Renen, S. Imamura, J. Giceva, T. Neumann, and A. Kemper, "Plush: A write-optimized persistent log-structured hashtable," *Proceedings of the VLDB Endowment*, vol. 15, no. 11, pp. 2895–2907, 2022.